



Neural Networks and Google TensorFlow with R!


SHANNON MCCORMICK

Outline

- ▶ Intro to Neural Networks
 - ▶ Inspiration
 - ▶ Basic explanation
 - ▶ Practical applications
- ▶ Google TensorFlow
 - ▶ What is it?
 - ▶ R Package
 - ▶ Basic Usage
 - ▶ Examples


Neural Network Inspiration

- ▶ Real Neurons
 - ▶ Dendrites receive an input
 - ▶ Based on input, axons output something to next neuron




Artificial Neural Networks

- ▶ Set of nodes connected by directional lines representing weights
 - ▶ Nodes represent mathematical operations
 - ▶ Weights learned by training




Linear Regression Example



- ▶ No hidden layers
- ▶ $\text{Inputs} * \text{Weights} = \text{Output}$
- ▶ Weights selected that minimize the error
- ▶ Great at modeling linear relationships

Adding Hidden Layers

- ▶ Add hidden layer(s)
- ▶ $\text{Input} \times \text{Weights}_1 = \text{Hidden Layer}$
- ▶ $\text{Hidden Layer} \times \text{Weights}_2 = \text{Output}$
- ▶ Weights selected to minimize error
- ▶ Can model more complex relationships



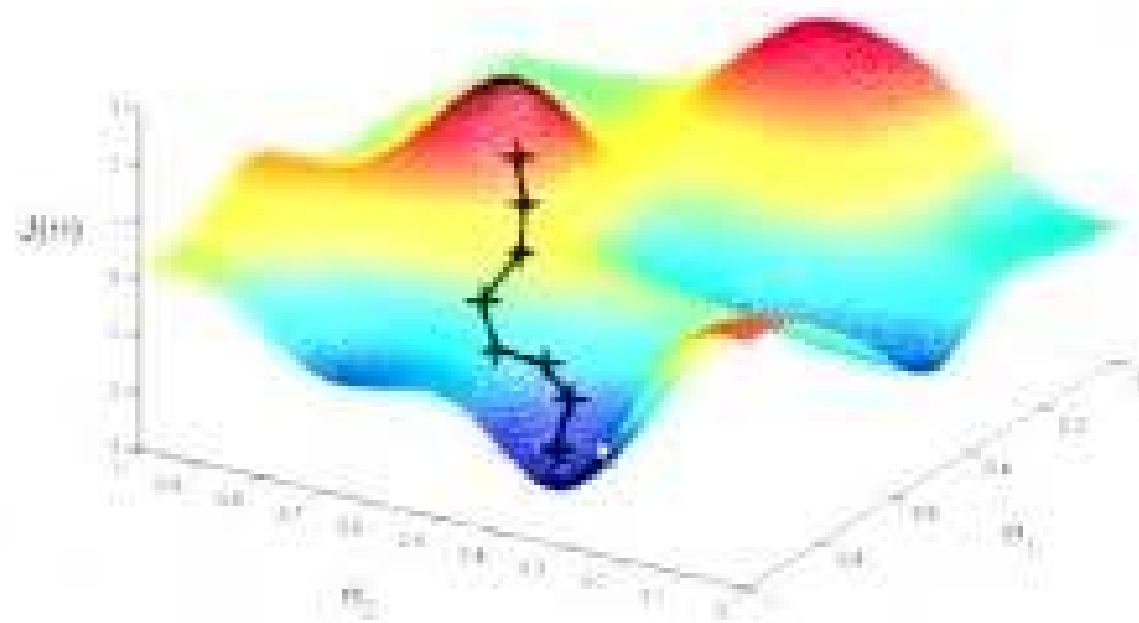
Learning

- ▶ Start
 - ▶ Weights have random initialization
 - ▶ Biases initialized at zero
- ▶ Forward Propagation
 - ▶ Batch of data goes through the network
 - ▶ Prediction/Classifications are output
- ▶ Error Calculation
 - ▶ Differences between true and outputted values are calculated

Learning weights


- ▶ Back Propagation
 - ▶ Errors are back propagated through the model
 - ▶ Determines the errors at each neuron in the network
- ▶ Gradient Descent
 - ▶ Optimization method
 - ▶ Determine how to change the weights
 - ▶ Takes a step down gradient of the function
- ▶ Iterative process

Gradient Descent



Google TensorFlow

- ▶ Open source machine learning library
- ▶ Released November 9, 2015
- ▶ Now the most popular machine learning framework



Features

- ▶ Can be used on desktop, mobile, servers
 - ▶ Linux, Mac OS, and Windows (added Nov, 2016)
 - ▶ GPU support on Linux and Mac OS (June, 2016)
- ▶ Written in C++ with Python interface
 - ▶ Added experimental Go (Nov, 2016) and Java (Feb, 2017) APIs
- ▶ Excellent step by step tutorials and documentation
- ▶ Auto-differentiation
- ▶ Includes Tensorboard for graph visualization
- ▶ Active improvement and growth
 - ▶ Google cloud (March 2016)
 - ▶ Distributed computing support (April 2016)
 - ▶ TensorFlow 1.0.0 released 5 days ago

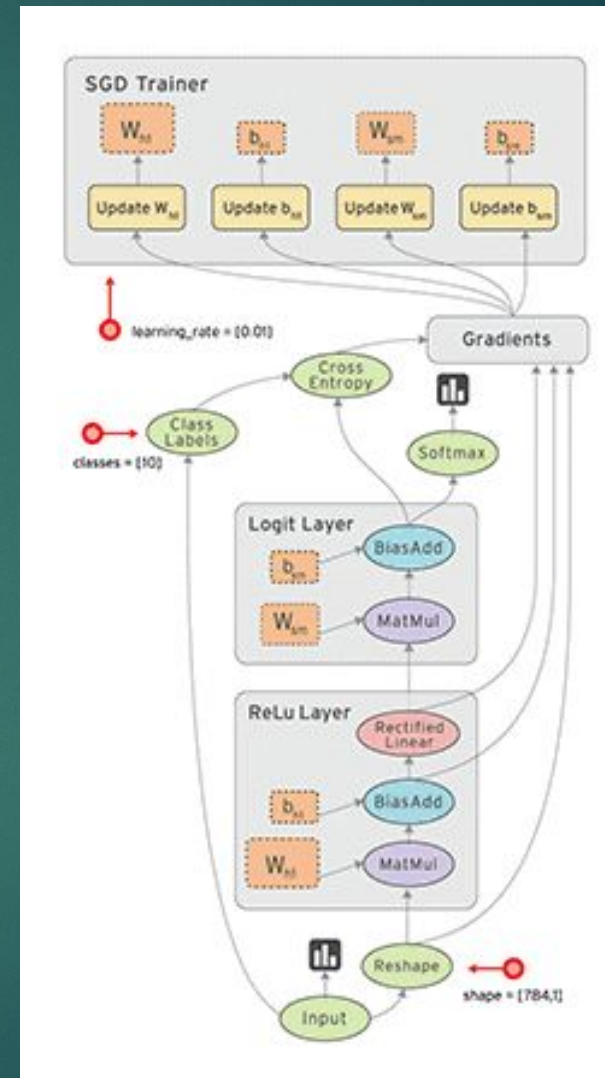
TensorFlow R Package

- ▶ Released October 2016
- ▶ Provides access to Python API through R
 - ▶ dependent on Python
 - ~~▶ Does not work with Anaconda~~
 - ~~▶ Not compatible with Windows~~
- ▶ When used with RStudio the package provides code completion and inline help
- ▶ Does include TFlearn higher level functions
- ▶ Cons
 - ▶ Not many examples
 - ▶ Points you to Python API documentation for

Basics

- ▶ Data flow graph with nodes and edges
 - ▶ Nodes: mathematical operations
 - ▶ Edges: input/output relationship between nodes
 - ▶ Edges carry tensors

Tensors flow through the graph



Building a model

“TensorFlow programs are usually structured into a construction phase, that assembles a graph, and an execution phase that uses a session to execute ops in the graph.” - TensorFlow docs


- ▶ Define computation graph
 - ▶ Inputs, operations, outputs
 - ▶ Symbolic representation of model
- ▶ Run session
 - ▶ Execute graph
 - ▶ Fetch output

Simplest Example

```
import tensorflow as tf

a = tf.constant(6.0)
b = tf.constant(7.0)
c = tf.mul(a, b)


with tf.Session() as sess:
    print(sess.run(c))
```



Simple Example

```
d = tf.placeholder(tf.float32)
e = tf.placeholder(tf.float32)
f = tf.add(a, b)

with tf.Session as sess:
    print(sess.run([f], feed_dict={d:[2.], e: [34.]}))
```



Classifying Covertypes

Covertypes data set from UCI Machine Learning Repository

- ▶ Predict 7 cover types from 54 input variables
 - ▶ 10 quantitative and 44 binary
- ▶ Data used for 2000 paper “Comparative Accuracies of Artificial Neural Networks and Discriminant Analysis in Predicting Forest Cover Types from Cartographic Variables.”
 - ▶ Trained neural network and attained 70.52% accuracy
 - ▶ Let's see if we can do better!

Some details

- ▶ 581,012 observation
 - ▶ Unbalanced categories
- ▶ Train, Test, and Validation
 - ▶ Train: 1620 observations of each covertime (11,340 obs)
 - ▶ 60% of observations of the least numerous cover type
 - ▶ Validation: 540 observations of each covertime (3,780 obs)
 - ▶ 20% of least numerous
 - ▶ Test: the rest (562,892 obs)
- ▶ Their network
 - ▶ 1 layer of 120 hidden nodes
 - ▶ Learning rate <- 0.05
 - ▶ Momentum rate <- 0.5

Additional Resources

- ▶ [TensorFlow Tutorials](#)
- ▶ [Udacity Deep Learning Course](#)
- ▶ [Awesome TensorFlow](#)
 - ▶ [TensorFlow Examples](#)
- ▶ [WildML](#)
- ▶ [TF Learn \(Scikit Flow\)](#)
- ▶ [Keras](#)
- ▶ [Stanford CS224d Lecture 7](#)
- ▶ [TensorBoard](#)
- ▶ [TensorFlow Playground](#)

